

UCL

Université
catholique
de Louvain



Computer Networking : Principles, Protocols and Practice

Part 2 : Applications

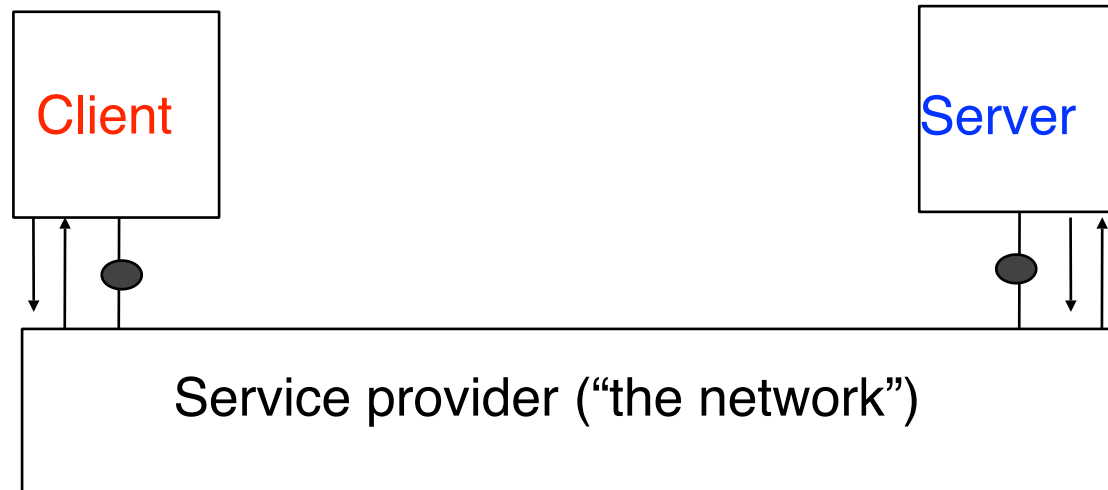
Olivier Bonaventure
<http://inl.info.ucl.ac.be/>

The Application Layer

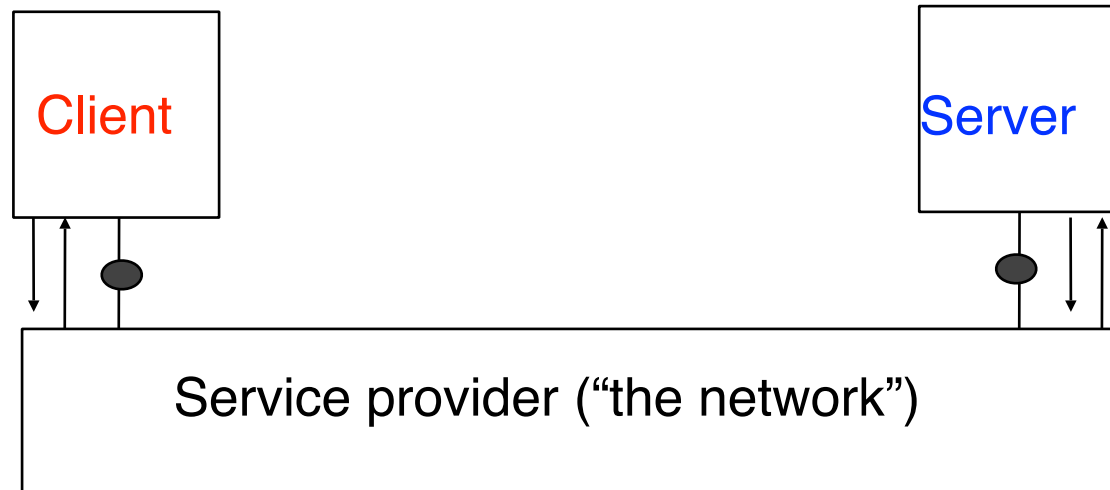
- Contents

- □ The client-server model
 - Name to address resolution
 - email
 - world wide web
 - peer-to-peer applications

The client server model

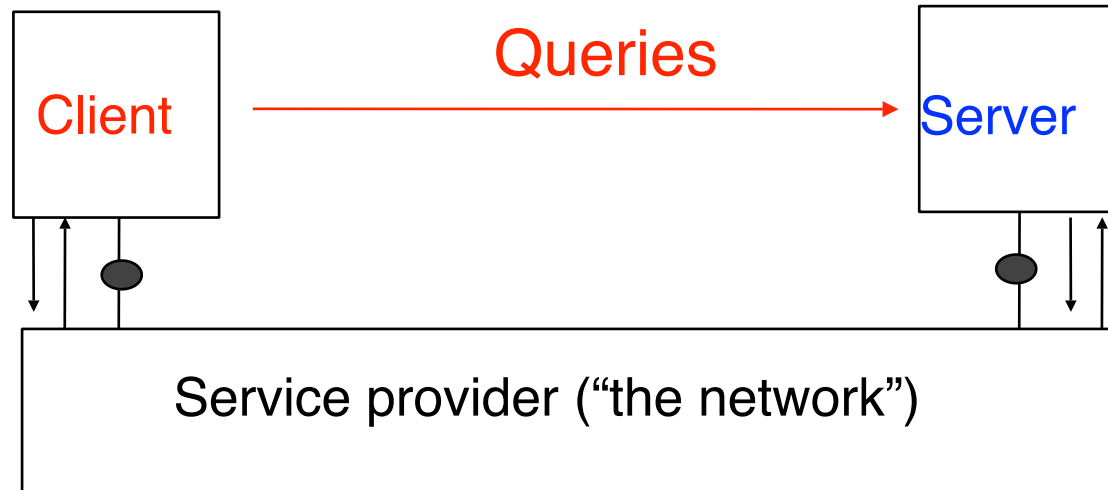


The client server model



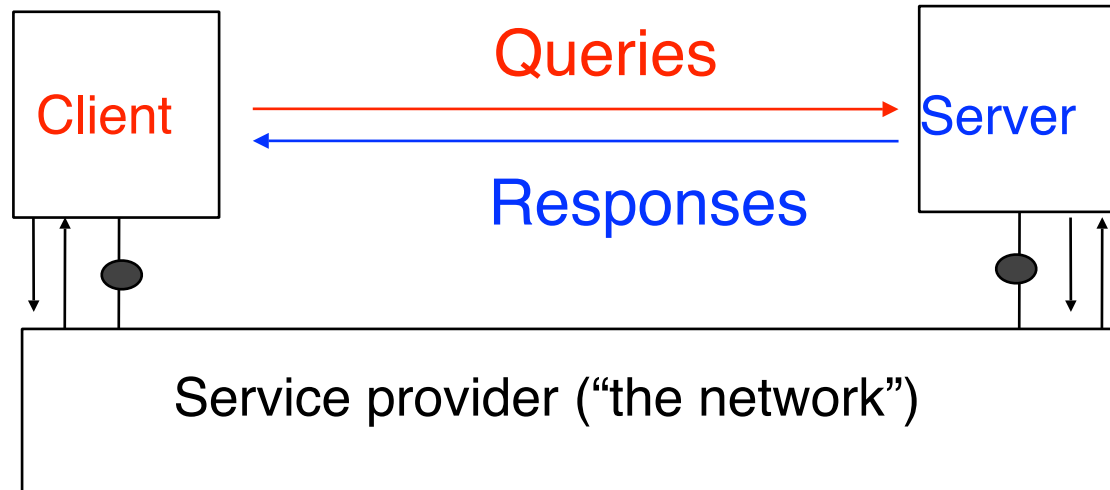
- Client
 - interacts with server through transport layer
 - sends queries or commands

The client server model



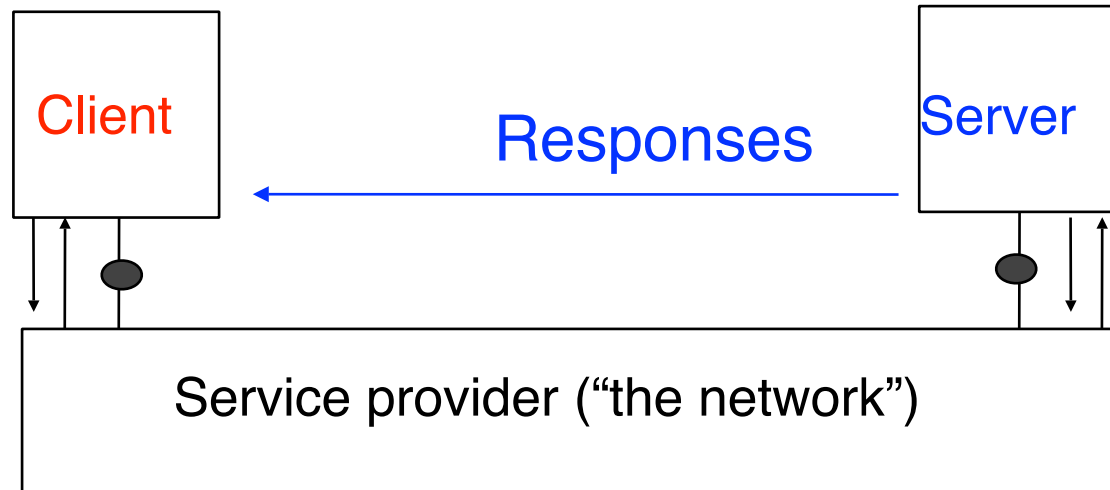
- Client
 - interacts with server through transport layer
 - sends queries or commands

The client server model



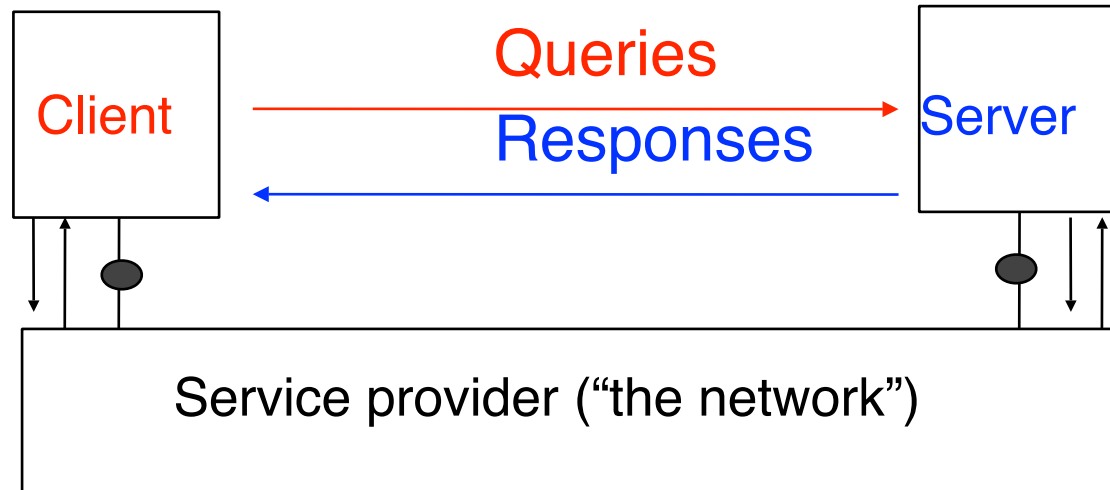
- Client
 - interacts with server through transport layer
 - sends queries or commands
- Server
 - Answers the queries received from clients
 - Executes the commands from clients
 - Many clients can use the same server
- Example : email, www, ...

The client server model (2)



- Client and servers interact with service provider
- Both the client and the server must speak the same language
 - Application-level protocol : set of syntactical and semantical rules that define the messages exchanged between the client and the server and their ordering

The client server model (2)

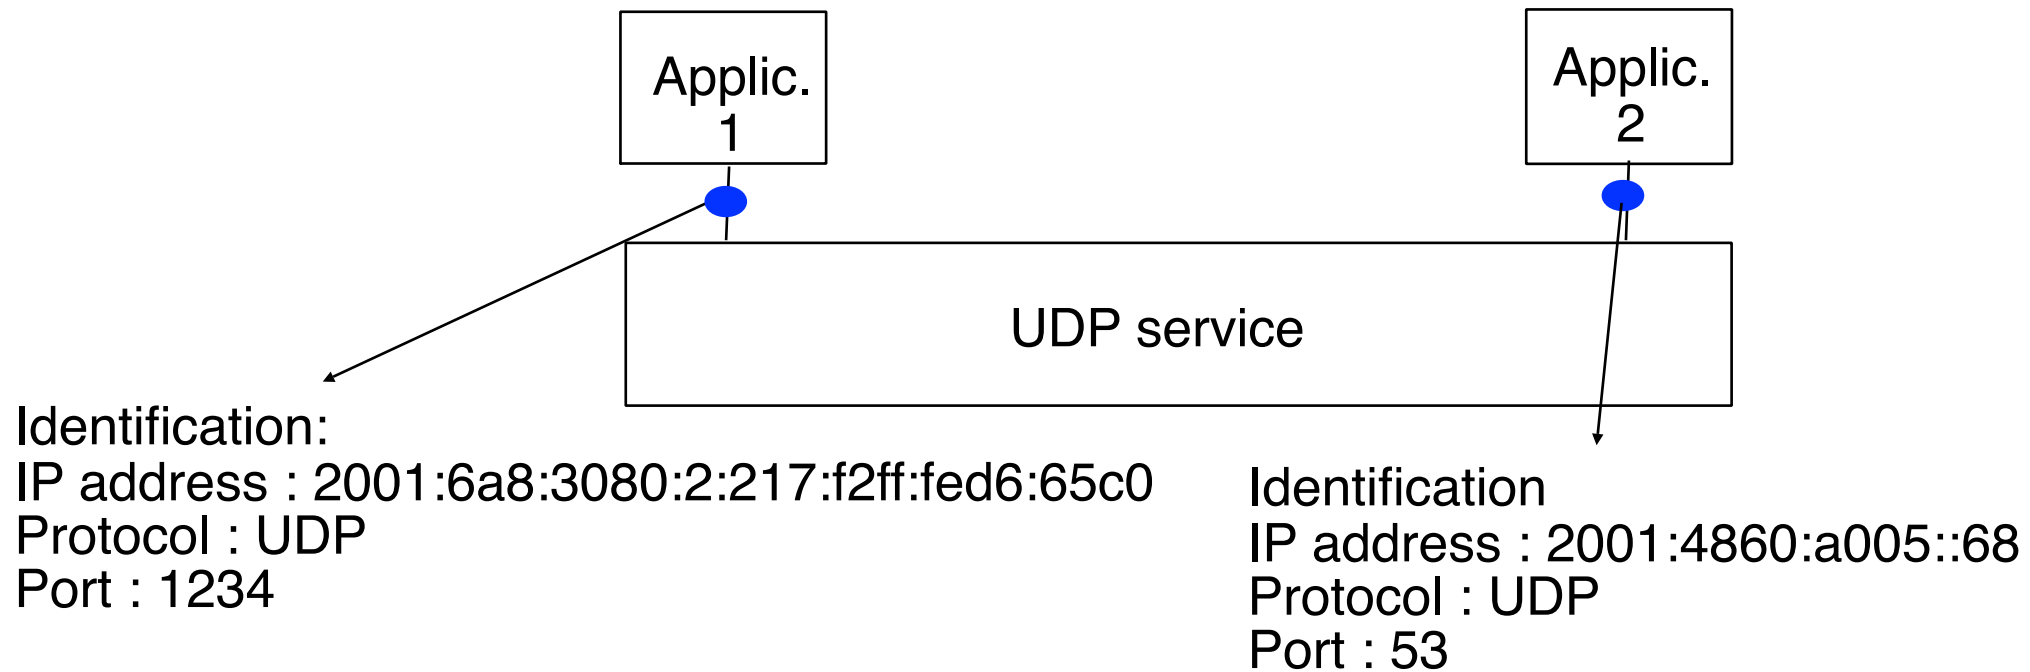


- Client and servers interact with service provider
- Both the client and the server must speak the same language
 - Application-level protocol : set of syntactical and semantical rules that define the messages exchanged between the client and the server and their ordering

Transport service on the Internet

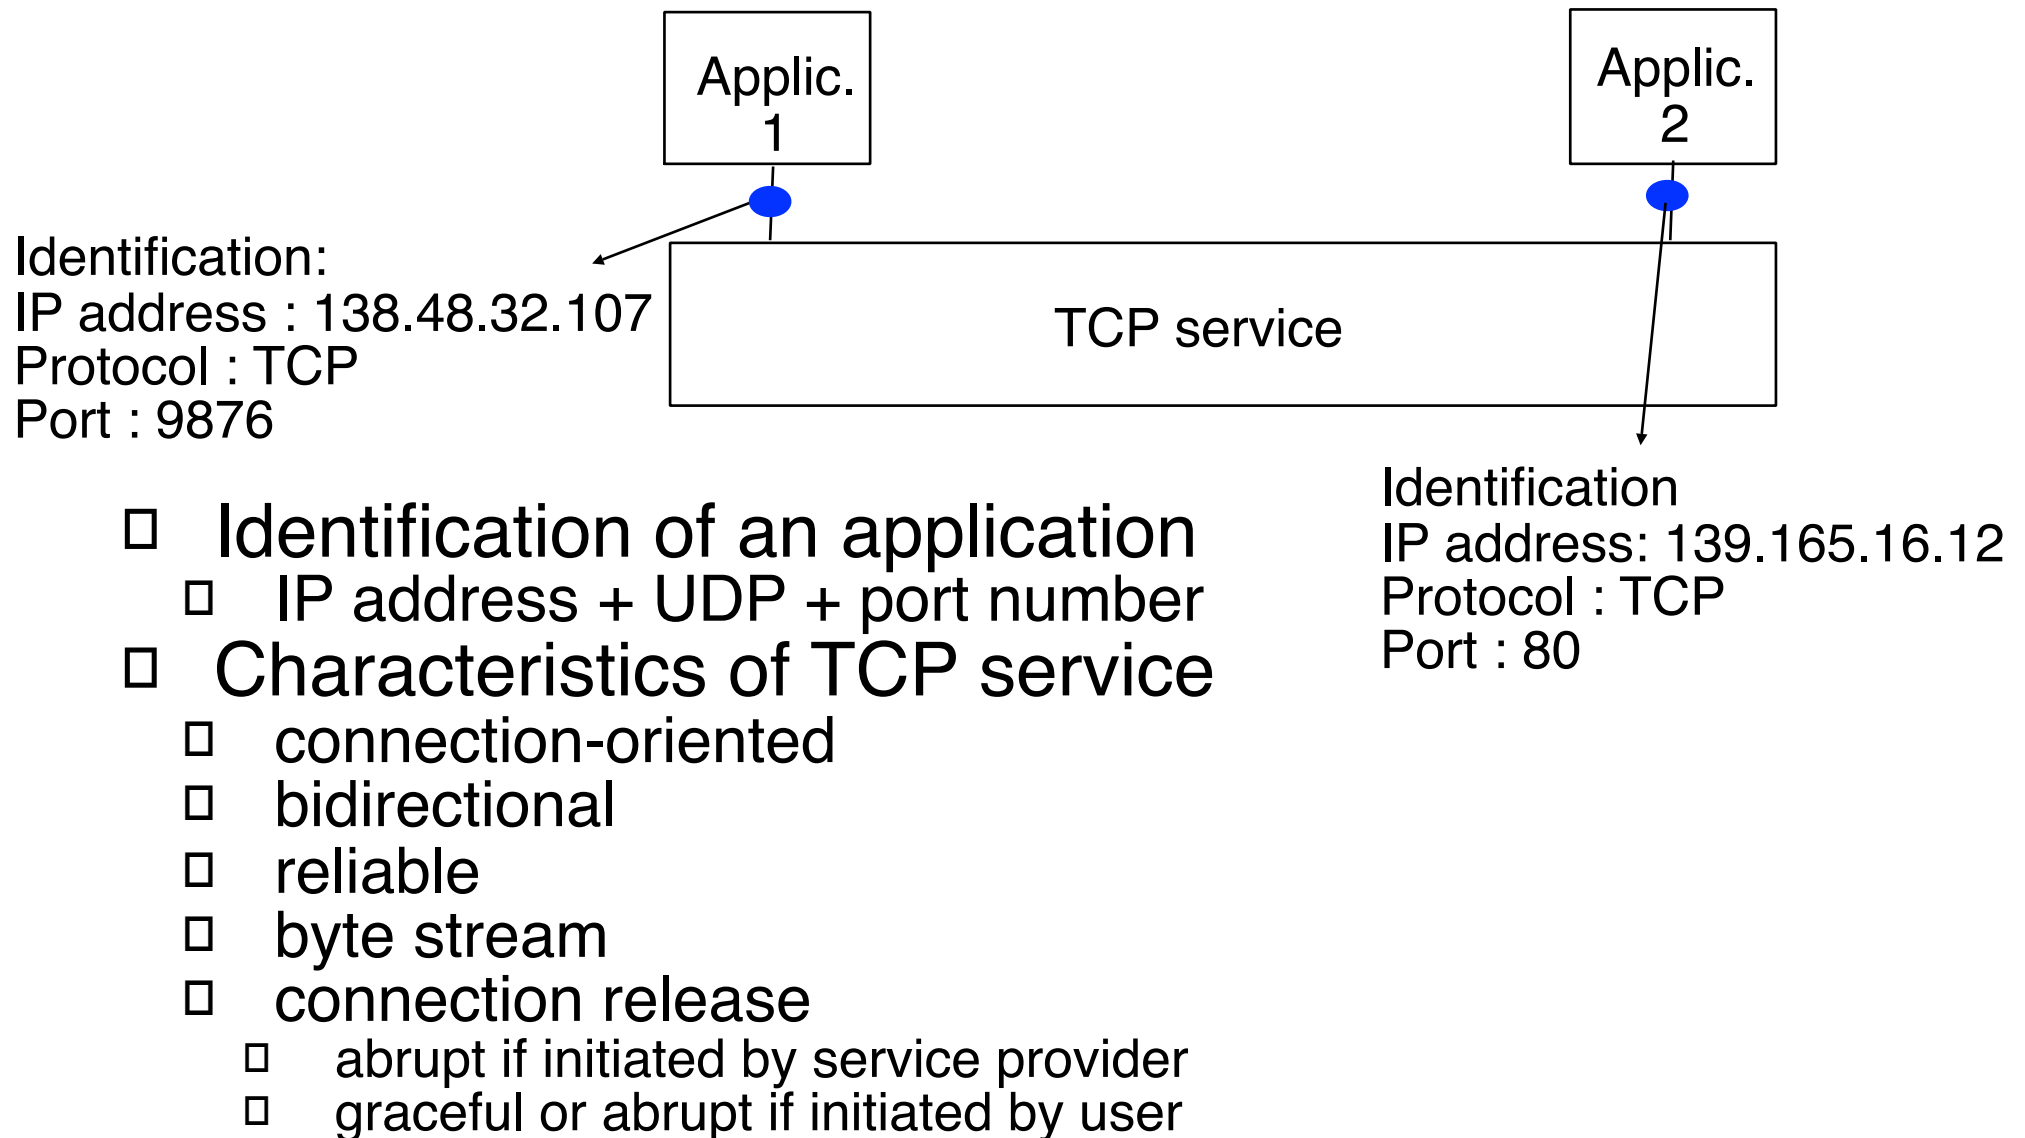
- On the Internet, applications can use two different transport services
 - The service provided by the User Datagram Protocol (UDP)
 - unreliable connectionless service with error detection
 - The service provided by the Transmission Control Protocol (TCP)
 - reliable bytestream connection-oriented service

UDP service



- ❑ Identification of an application
 - ❑ IP address + UDP + port number
- ❑ Characteristics of UDP service
 - ❑ connectionless
 - ❑ unreliable
 - ❑ messages can be lost
 - ❑ transmission errors can be detected but not recovered
 - ❑ sequence is not preserved
 - ❑ Maximum size of messages : almost 64 Kbytes

TCP service



Internet applications

- Contents

- The client-server model

- □ Name to address resolution

- email

- world wide web

- peer-to-peer applications

Names and addresses

Names and addresses

- Address of a server
 - IP Address of the host on which the server is running port number (TCP or UDP)
 - usually well known port number

Names and addresses

- Address of a server
 - IP Address of the host on which the server is running port number (TCP or UDP)
 - usually well known port number
- Drawback
 - Difficult to remember an IP address for a human

Names and addresses

- Address of a server
 - IP Address of the host on which the server is running
 - port number (TCP or UDP)
 - usually well known port number
- Drawback
 - Difficult to remember an IP address for a human
- Idea
 - Replace IP address by a hostname
 - Easier for humans
 - but IP address is necessary to contact server
- How to translate a hostname in an IP address ?

Names and addresses (2)

- ❑ `hosts.txt` file
 - ❑ contains the name-address table
 - ❑ must be updated regularly

```
#  
# Internet host table  
#  
127.0.0.1      localhost  
138.48.32.99   babbage  
138.48.32.100  leibniz  
138.48.32.1    routeur  
138.48.32.92   corneille  
138.48.32.107  backus  
138.48.20.152  arzach  
138.48.32.137  almin01  
138.48.32.170  duke
```

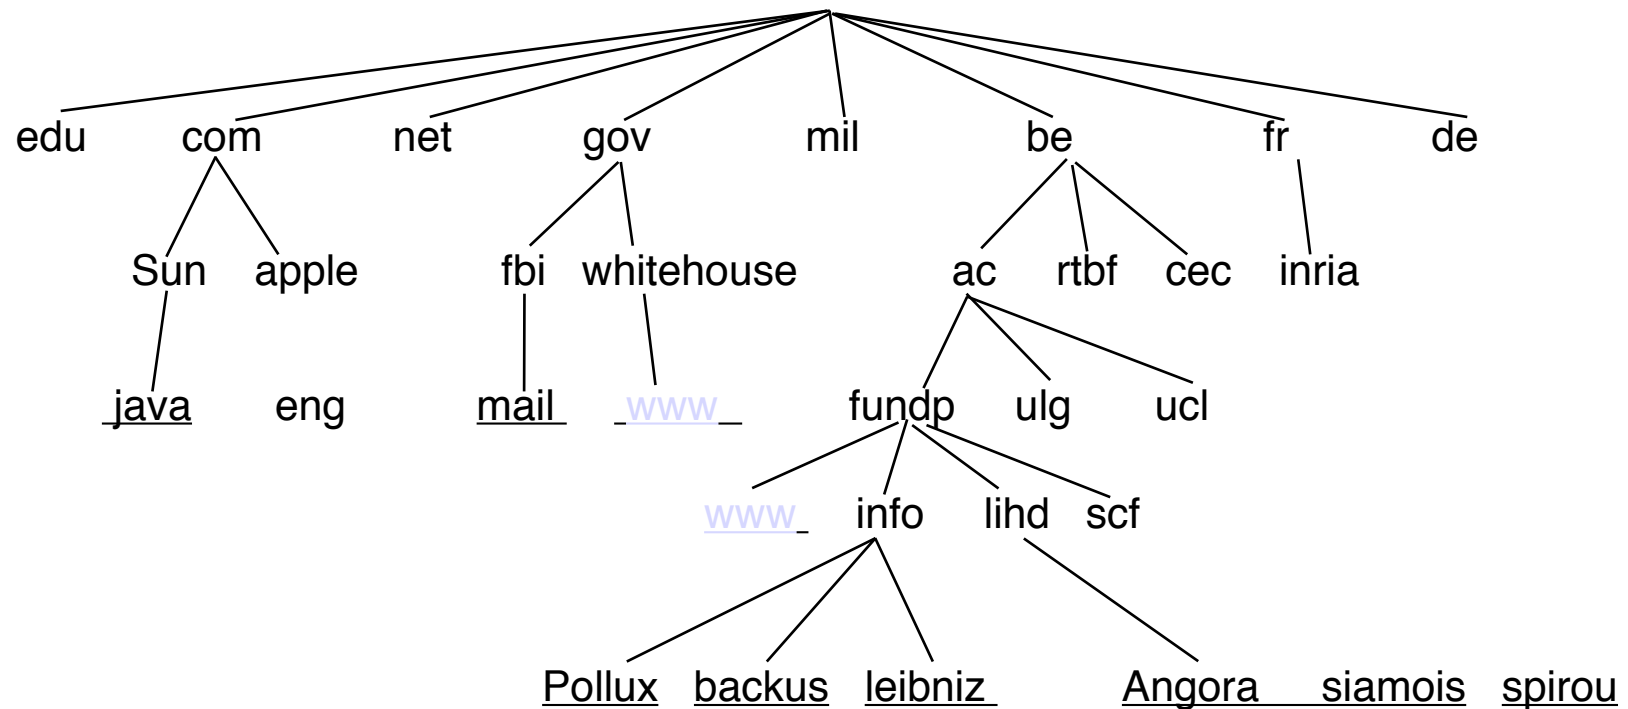
- ❑ cannot be used in a large network

Hostnames

- Requirement
 - Host names should be unique
- How to achieve this in a scalable manner ?
 - Introduce hierarchy
 - Each hostname is composed of two parts
 - domain name (globally unique)
 - hostname (unique within a given domain)
- How to uniquely distribute domain names ?
 - Introduce hierarchy
 - A small number of top-level domain names
 - Inside each top-level domain, allocate uniquely second level domain names
 - Inside each seconde-level domain, allocate uniquely either third-level domain names or host names,
 - ...

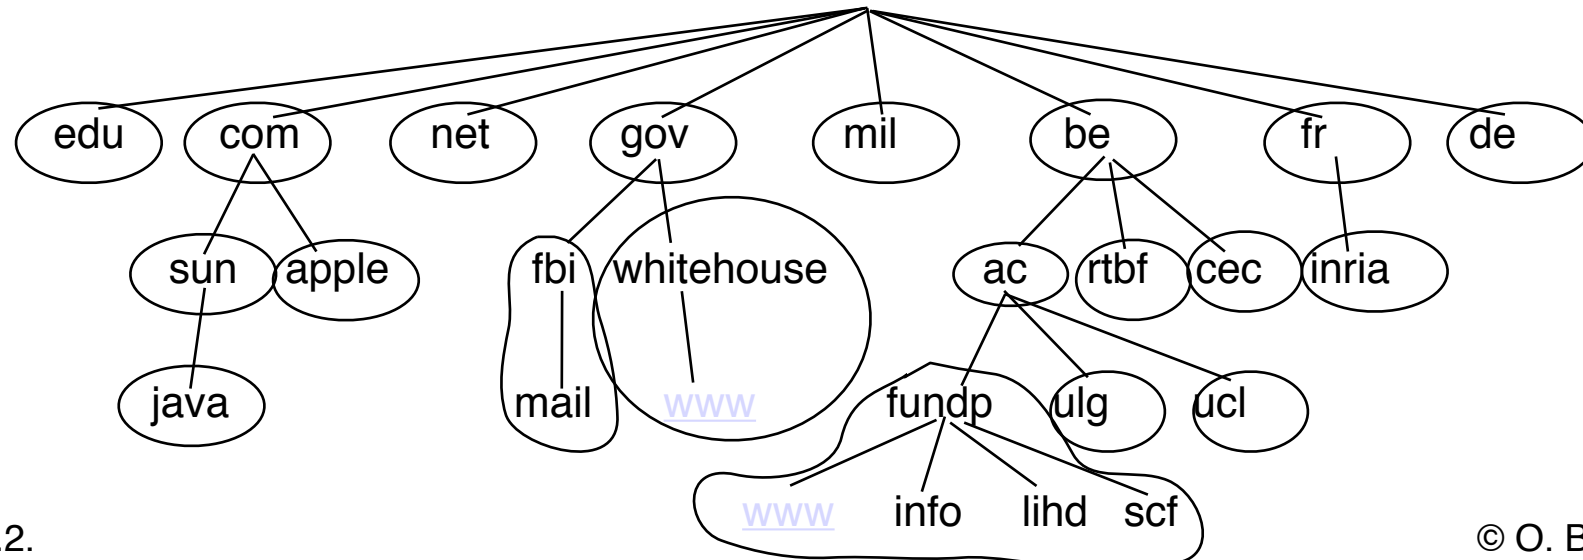
Host names and domain names

□ Tree of all host names



How to translate names into addresses ?


- How to efficiently translate a host name ?
 - By using a centralised database
 - there are more than 1 billion host names today
 - By using a distributed database
 - DNS : Domain Name System
 - relies on the hierarchy of domain names
 - there is one server responsible for each domain and this server must be queried to translate host names inside this domain



How to translate names into addresses ?

- Domain Name Service (DNS)
 - Each DNS server is responsible for a domain and knows
 - The IP addresses of all host names in this domain
 - The IP addresses of the DNS servers responsible for subdomains

□ Examples



- java.sun.com
- www.ucl.ac.be

DNS resolver

- ❑ To be able to translate name to addresses, a DNS implementation needs
 - ❑ to know **actual** list of IP addresses of root servers
 - ❑ to implement the DNS protocol and traverse the domain names hierarchy
 - ❑ Difficult to do this on all endhosts

- ❑ Solution
 - ❑ DNS resolver
 - ❑ one resolver for a set of endhosts
 - ❑ maintains up-to-date list of IP addresses of root servers
 - ❑ implements DNS protocol

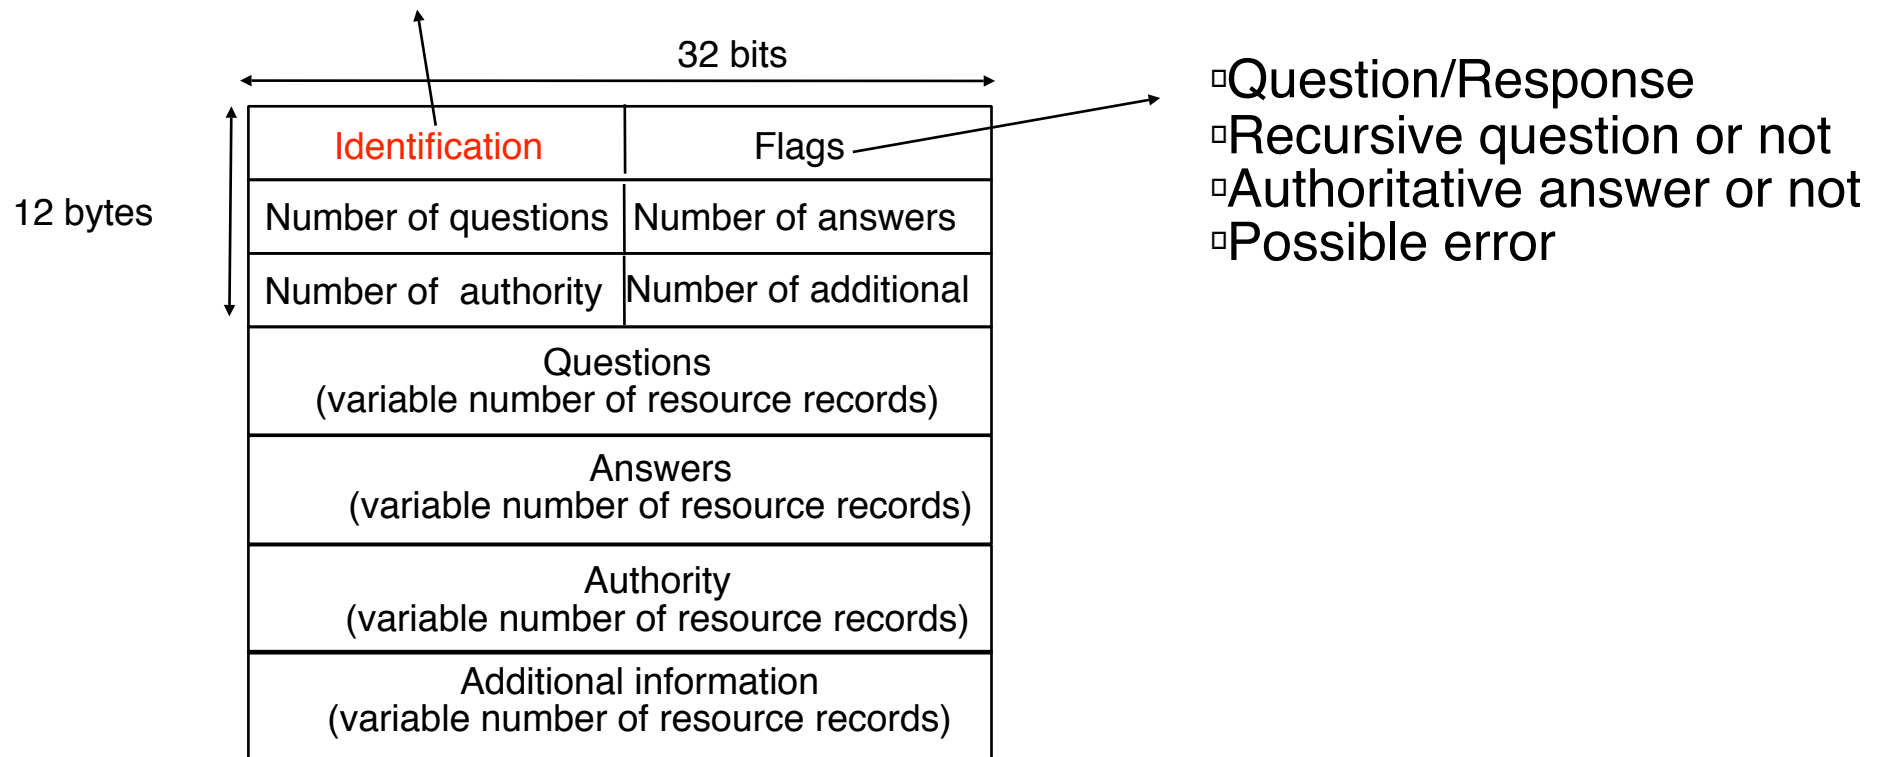
 - ❑ endhosts
 - ❑ only need to be able to send DNS requests to resolver
 - ❑ must know IP address of closest DNS resolvers

DNS : optimisations

- ❑ Reduce risk of failures
 - ❑ several root-servers
 - ❑ server DNS servers authoritative for each domain
 - ❑ each endhost can send queries to multiple resolvers
- ❑ Improved performance
 - ❑ avoid sending several times the same query
 - ❑ cache memory on DNS resolvers containing
 - ❑ recent name-addresses translations
 - ❑ addresses of DNS servers recently contacted
- ❑ DNS protocol
 - ❑ usually runs over UDP
 - ❑ sometimes is also used over TCP

DNS : message format

Each DNS request contains a number that will be returned in the response by the server to allow the client to match the request.



DNS : resource records

- Each DNS messages is composed of resource records (RR) encoded as **TLV**
 - **< Name, Value, Type, TTL >**
 - Types de RR
 - A (Address)
 - **Name** is a hostname and **Value** an **IPv4 address**
 - AAAA (Address)
 - **Name** is a hostname and **Value** an **IPv6 address**
 - NS (NameServer)
 - **Name** is a domain name and **Value** is the hostname of the DNS server responsible for this domain
 - MX (Mail Exchange)
 - **Name** is a domain name and **Value** is the name of the SMTP server that must be contacted to send emails to this domain
 - Type CNAME
 - Alias
- ↙ Lifetime of the RR in server's cache

Internet applications

- Contents

- The client-server model

- Name to address resolution

- □ email

- world wide web

- peer-to-peer applications

Email

- Simplified model
 - Alice sends an email to Bob



Alice@a.net



Alice's
email server



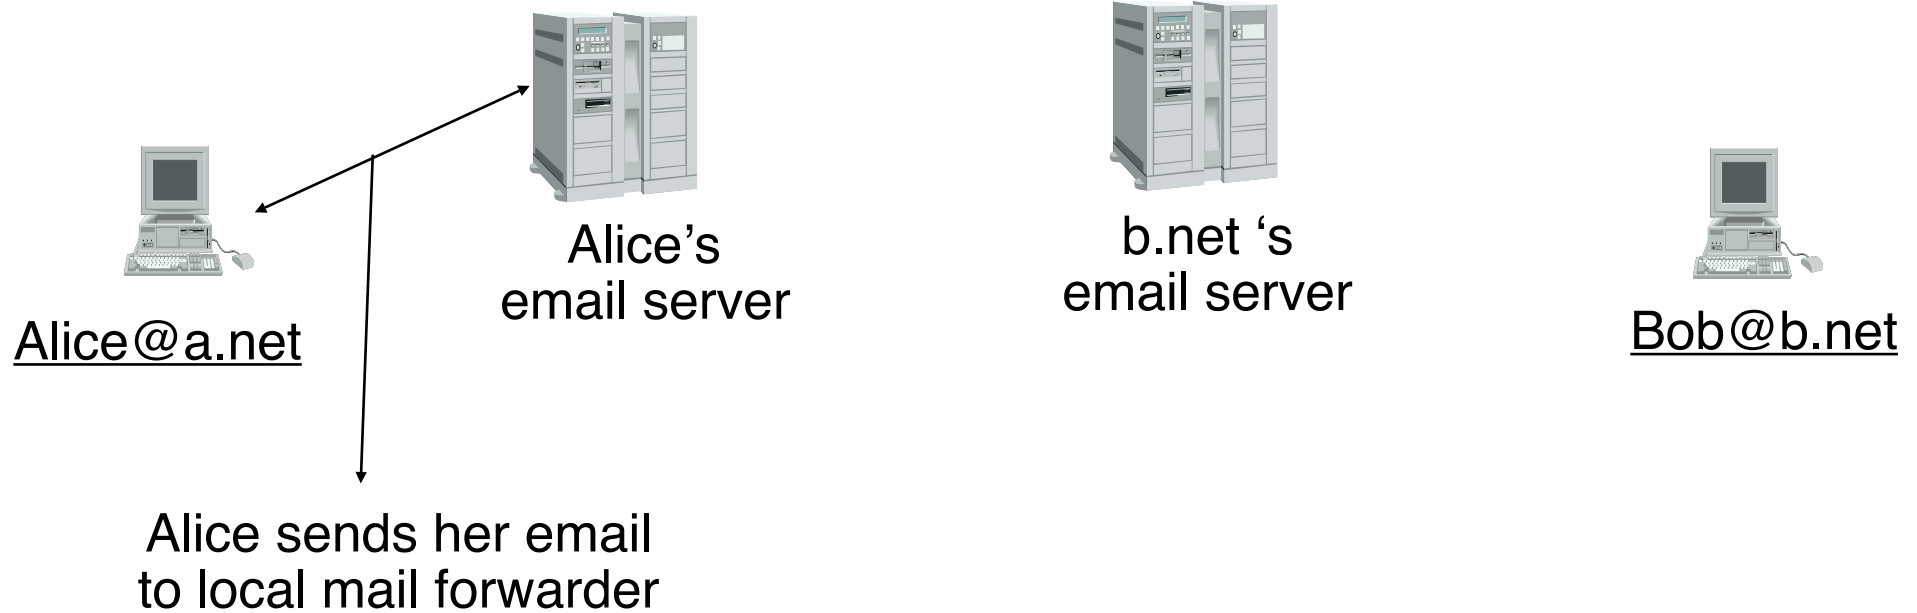
b.net 's
email server



Bob@b.net

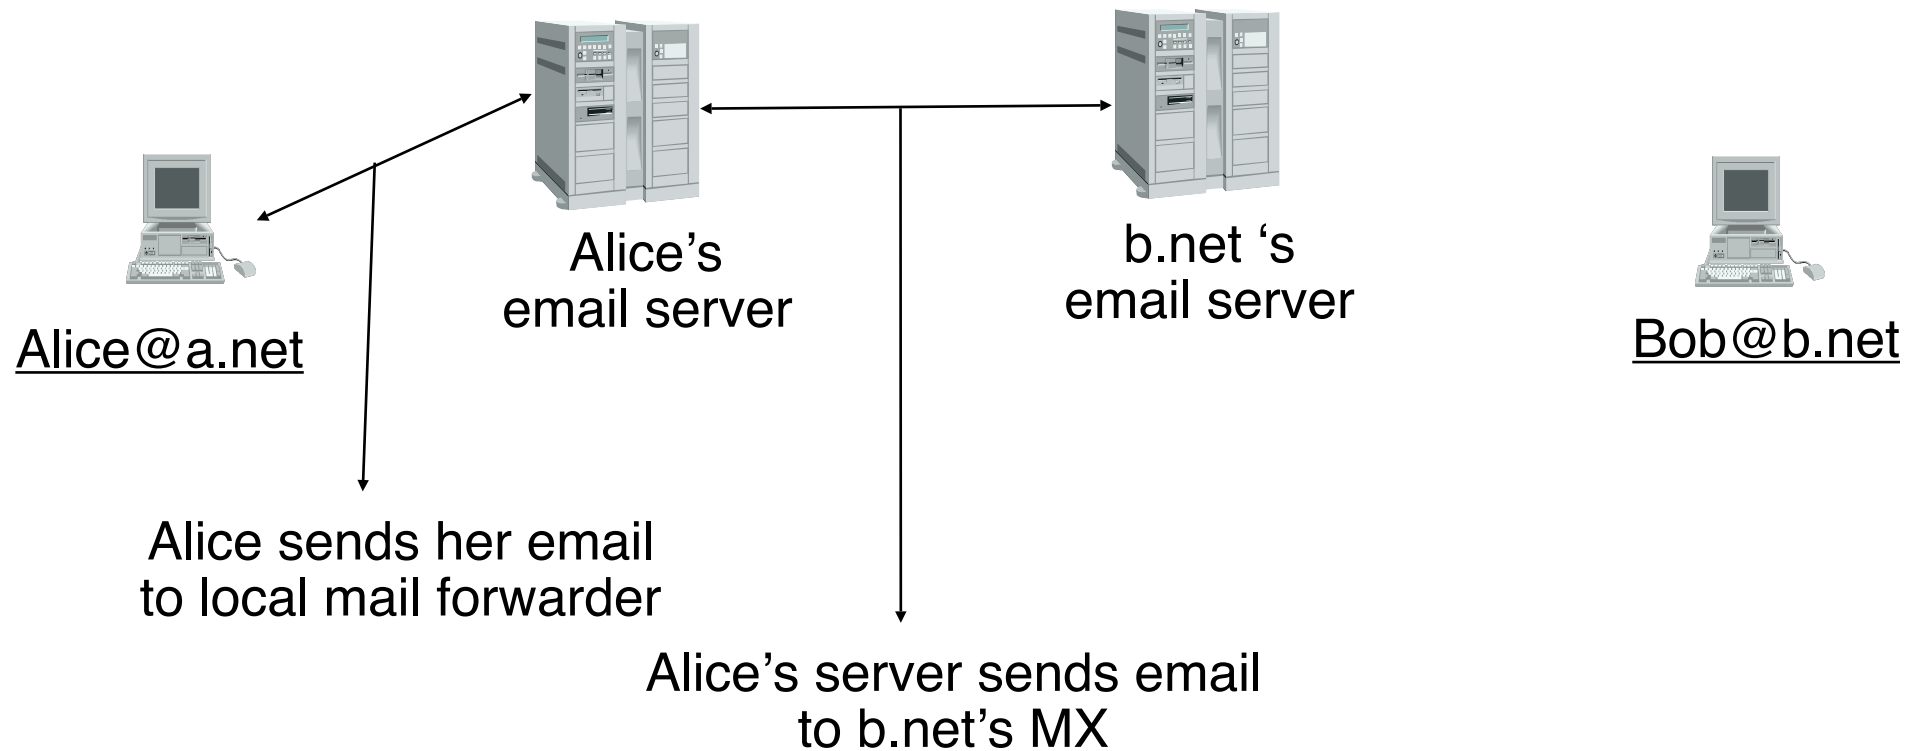
Email

- Simplified model
 - Alice sends an email to Bob



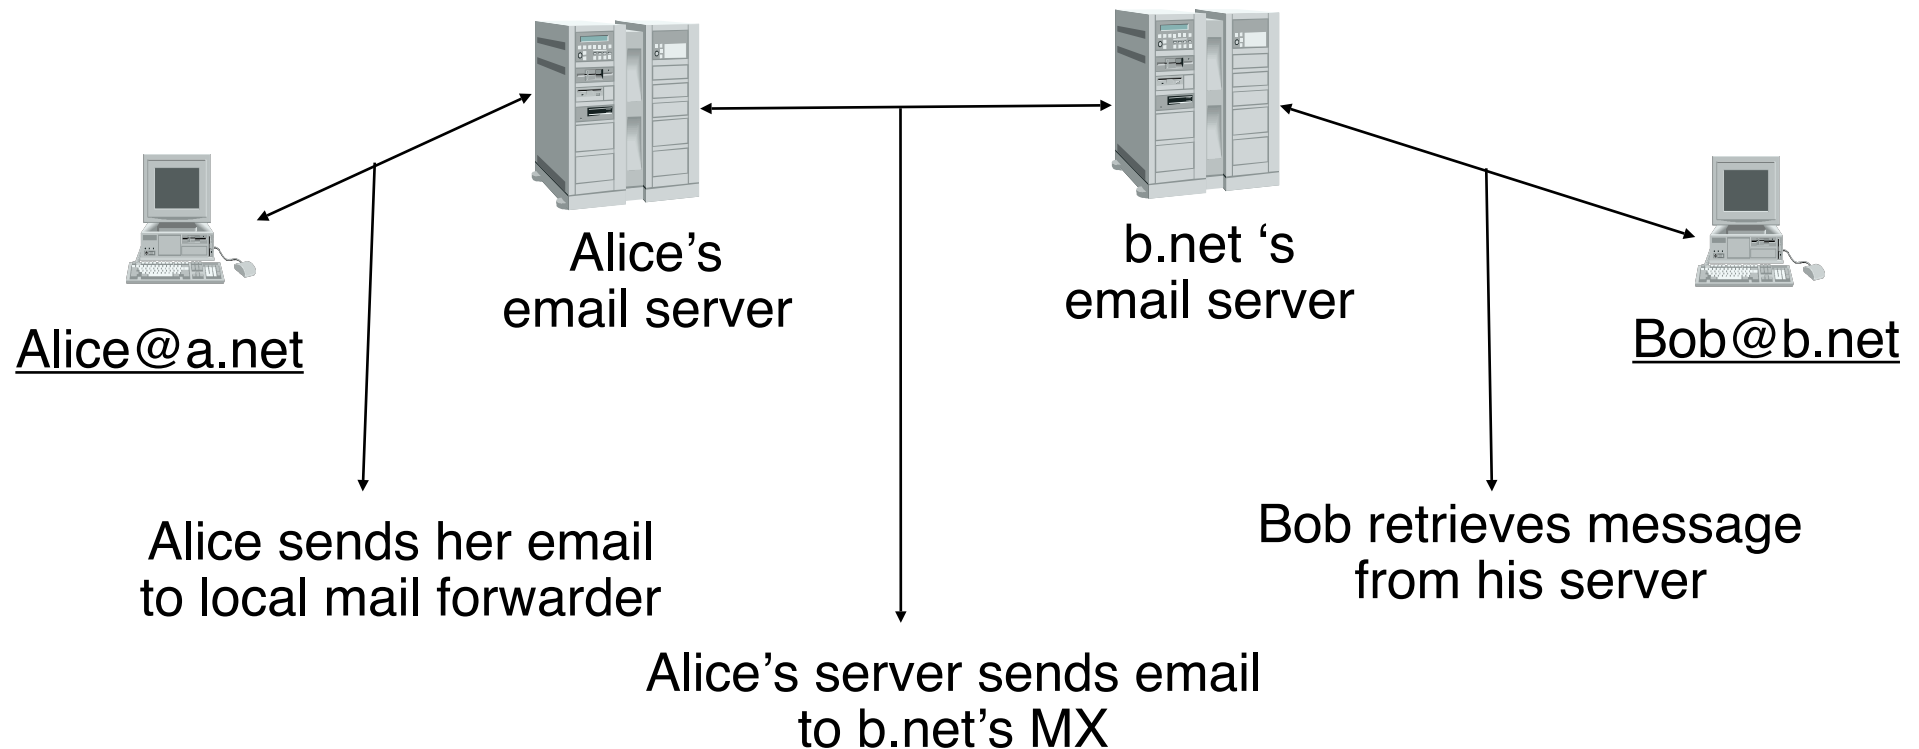
Email

- Simplified model
 - Alice sends an email to Bob

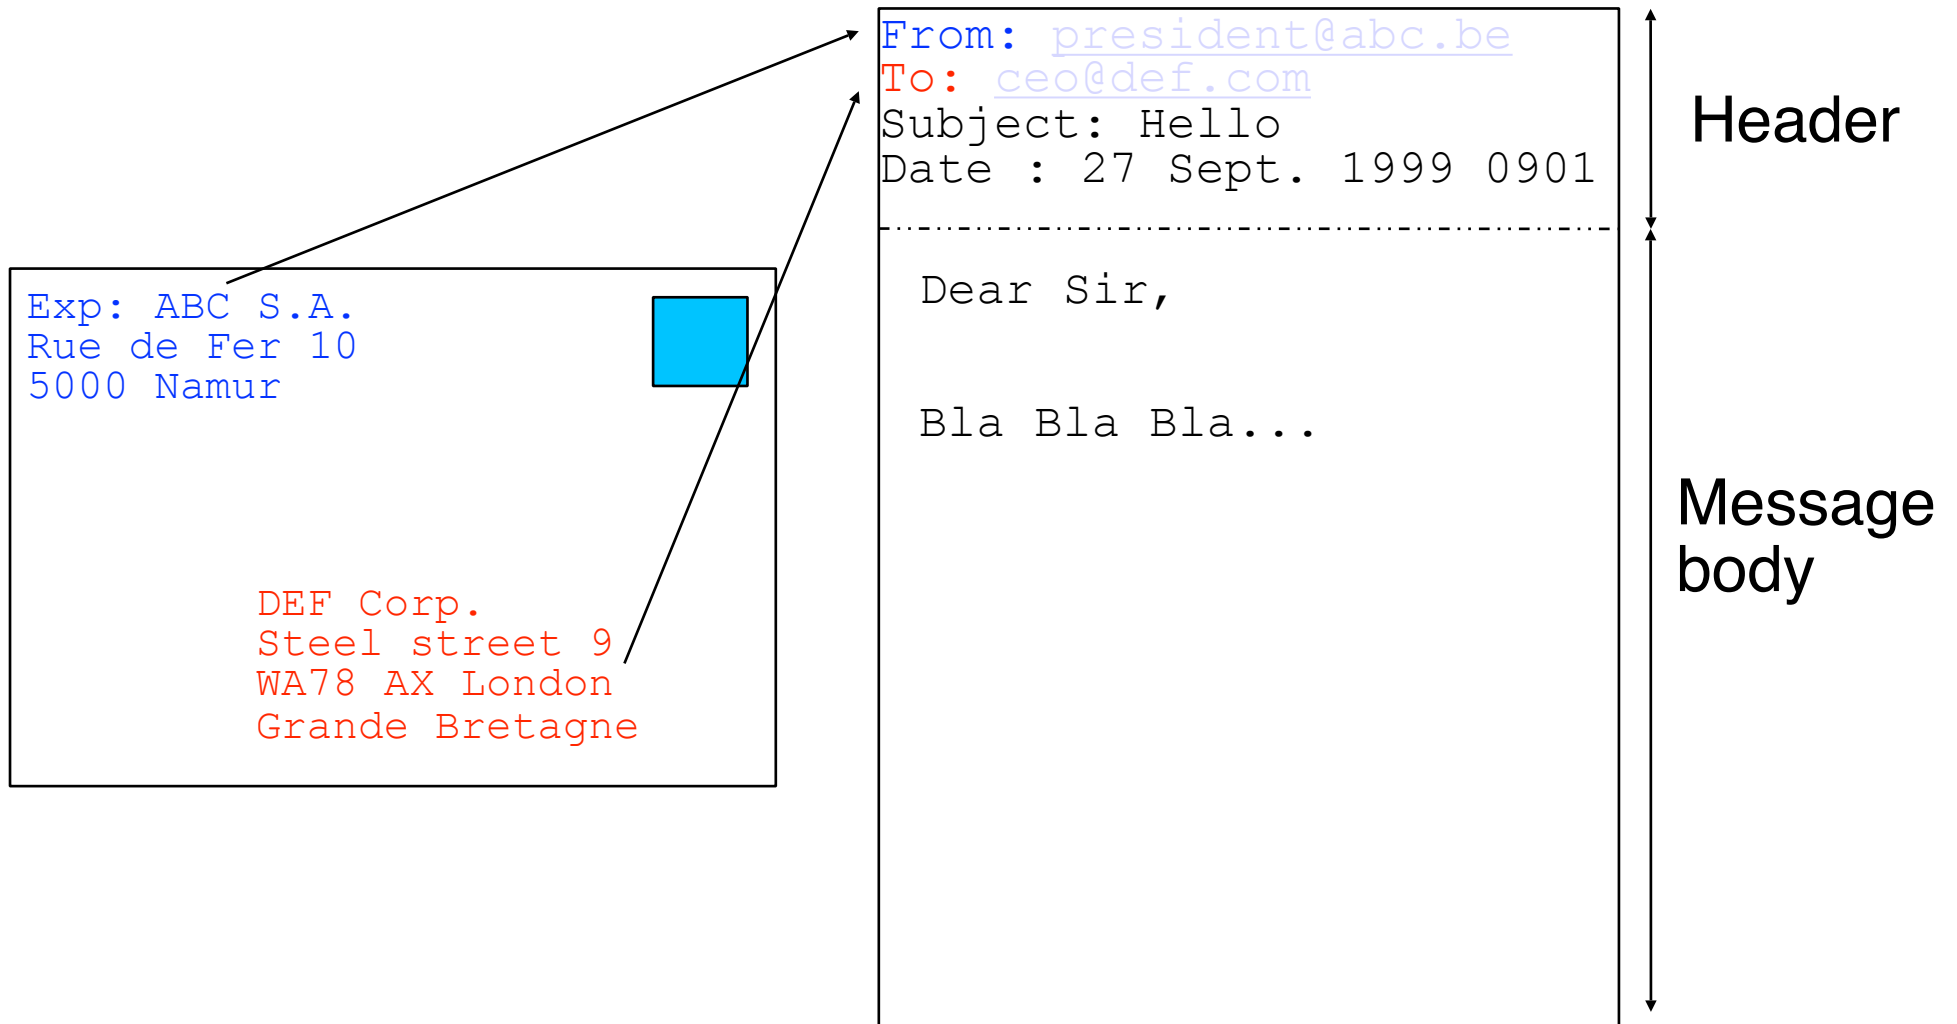


Email

- Simplified model
 - Alice sends an email to Bob



Email message format



Message format (2)

- ❑ Header format
 - ❑ Contains only US-ASCII (7bits) characters
 - ❑ At least three lines that end with <CRLF>
 - ❑ From: sender@domain
 - ❑ To: [recipient@domain](#)
 - ❑ Date: <creation date of message>
 - ❑ example : 26 Aug 199 1445 EDT
 - ❑ Optional fields
 - ❑ Subject: subject of message
 - ❑ cc: [copy@domain](#)
 - ❑ Message-ID: <[number@domain](#)>
 - ❑ Received: information on path followed by message
 - ❑ In-Reply-To: <message-ID>
 - ❑ Header ends with empty line (<CRLF>)

MIME

MIME

- Internet email was designed for US-ASCII
 - How to transmit more complex messages ?

MIME

- ❑ Internet email was designed for US-ASCII
 - ❑ How to transmit more complex messages ?
- ❑ Multipurpose Internet Mail Extensions
 - ❑ Improved email message format
 - ❑ Constraints
 - ❑ must remain compatible with old email servers
 - ❑ most of them only support US-ASCII and short lines
 - ❑ must support non-English text
 - ❑ character set must be beyond 7bits US-ASCII
 - ❑ must support various formats in a single message
 - ❑ message body, attachments, ...
 - ❑ must allow to transmit audio, video, ...
 - ❑ need to identify the type of content

MIME

- ❑ Internet email was designed for US-ASCII
 - ❑ How to transmit more complex messages ?
- ❑ Multipurpose Internet Mail Extensions
 - ❑ Improved email message format
 - ❑ Constraints
 - ❑ must remain compatible with old email servers
 - ❑ most of them only support US-ASCII and short lines
 - ❑ must support non-English text
 - ❑ character set must be beyond 7bits US-ASCII
 - ❑ must support various formats in a single message
 - ❑ message body, attachments, ...
 - ❑ must allow to transmit audio, video, ...
 - ❑ need to identify the type of content
- ❑ Solution
 - ❑ add new optional fields in header
 - ❑ add optional fields inside message body when

MIME (2)

□ New header fields

□ MIME-Version:

- version of MIME used to encode message
- current version : 1.0

□ Content-Description:

- comment describing the content of the message

□ Content-Type:

- type of information inside message

□ Content-Transfer-Encoding:

- how the message has been encoded

□ Content-Id:

- unique identifier for the content

MIME: Content-Type

- ❑ Content-Type : type/encoding
 - ❑ type of content
 - ❑ text, image, video, application
 - ❑ multipart
 - ❑ encoding of content
 - ❑ text/plain , text/html
 - ❑ image/gif, image/jpeg
 - ❑ audio/basic
 - ❑ video/mpeg, video/quicktime
 - ❑ application/octet-stream, application/postscript
 - ❑ multipart/alternative
 - ❑ message contains several times the same information with different encodings
 - ❑ multipart/mixed
 - ❑ message contains several information of different types
 - ❑ example : text of message body and attachment

Character sets and content encoding

- ❑ How to support rich character sets ?
 - ❑ Content-Type: text/plain; charset=us-ascii
 - ❑ ASCII 7bits, default
 - ❑ Content-Type: text/plain; charset=iso-8859-1
 - ❑ Character set suitable for Western European languages, defined by ISO, 8 bits per character
 - ❑ Content-Type: text/plain; charset=unicode
 - ❑ Universal character set, defined by ISO, 16 bits per character
- ❑ How to encode non-text data ?
 - ❑ data must be encoded in US-ASCII 7 bits characters
 - ❑ Base64
 - ❑ uses ASCII caracteres A...Z,a...z,0...9, "+" et "/"
 - ❑ A=0, B=1, C=2, ... +=62 et /=63
 - ❑ Each character is used to encode 6 bits
 - ❑ 24 bits from initial message -> 4 ASCII characters
 - ❑ Special character "=" used for padding

Multipart/mixed

- How to place different contents and encoding in a single message ?
 - We need a delimiter between the different content types placed inside message body

```
Date: Mon, 20 Sep 1999 16:33:16 +0200
From: Nathaniel Borenstein <nsb@bellcore.com>
To: Ned Freed <ned@innosoft.com>
Subject: Test
MIME-Version: 1.0
Content-Type: multipart/mixed; boundary="simple boundary"

preamble, to be ignored

--simple boundary
Content-Type: text/plain; charset=us-ascii


partie 1

--simple boundary
Content-Type: text/plain; charset=us-ascii

partie 2
--simple boundary
```


Email transmission

- SMTP : Simple Mail Transfer protocol
 - uses TCP service
 - Address of SMTP server
 - IP address of server + TCP + port number: 25
 - RR of type MX can be used to find the SMTP server responsible for a given domain



SMTP

- Client-server model
 - Server waits for email messages to relay/deliver
 - Client sends email messages through server

- Application-level protocol
 - client opens TCP connection
 - Client sends commands composed of
 - command parameter <CRLF>
 - HELO
 - MAIL FROM:
 - RCPT TO:
 - DATA
 - QUIT
 - Server answers with one-line replies
 - numeric_code comment (text) <CRLF>
 - 250 OK
 - 221 closing

SMTP (2)

- Three phases of SMTP
 1. Establishment of an SMTP association
 - TCP connection established upon request from client
 - Server greetings
 - HELO command from client
 2. Message transfer
 - MAIL FROM: <[user@domaine](#)>
 - RCPT TO: <[user@domaine](#)>
 - DATA
 - transmission of entire message including headers
 - one line containing only the dot “.” characters marks end of message
 - Other subsequent messages can be transmitted after
 3. Release of the SMTP association
 - QUIT
 - Closing message from server
 - TCP connection is closed

Retrieval of email messages

□ In the old days

1. Destination is always connected to the Internet

- email addresses are username@hostname
- When an email arrives, it is stored in a file that belongs to the user, e.g. `/var/mail` on Unix

□ Today

- Most networks have one or a few SMTP servers used to receive emails, but also detect spam, viruses, ...
- Endusers retrieve their emails from this server
 - Post Office Protocol (POP)
 - Internet Mail Access Protocol (IMAP)
 - Webmail

POP

- Goal
 - Allow authenticated users to retrieve email messages from server
- Operation
 - POP uses TCP service
 - Address of POP server
 - Host address + TCP + port number : 110
 - Client send commands
 - command : one ASCII line ending with <CRLF>
 - USER, PASS, STAT, RETR, DELE, QUIT
 - server replies with
 - +OK if command was successful
 - email messages follow some +OK replies
 - -ERR in case of errors

POP (2)

- ❑ Three phases of the protocol
 1. Authorisation : checking the user credentials
 - ❑ USER <username>
 - ❑ PASS <password>
 2. Transaction
 - ❑ retrieval and removal of messages
 - ❑ STAT
 - ❑ list headers of stored messages
 - ❑ RETR <n>
 - ❑ retrieval of the nth message
 - ❑ DELE <n>
 - ❑ the nth message is marked for deletion
 3. Update
 - ❑ End of the retrieval phase
 - ❑ Messages marked for deletion are removed from server
 - ❑ TCP connection is closed

Internet applications

- Contents

- The client-server model

- Name to address resolution

- email

- □ world wide web

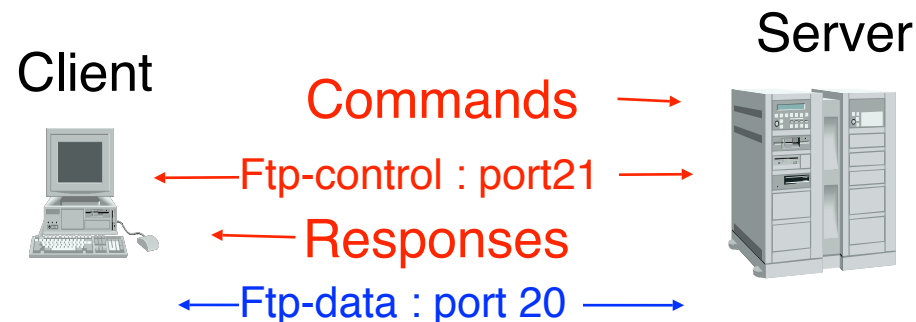
- peer-to-peer applications

FTP : File Transfer Protocol

- Protocol from the old days
 - allows a client to send/retrieve files from a server
- Problems solved by FTP
 - User authentication
 - username, password
 - Filesystem traversal
 - browse directories on server and locate files
 - file transfer
 - to or from the server

FTP : File Transfer Protocol

- Protocol from the old days
 - allows a client to send/retrieve files from a server
- Problems solved by FTP
 - User authentication
 - username, password
 - Filesystem traversal
 - browse directories on server and locate files
 - file transfer
 - to or from the server




FTP : main commands

□ Main commands

- USER <user>
 - **username**, ftp for anonymous access
- PASS <pass>
 - allows user to send **password** associated to username
- SYST
 - information about type of server
- CWD <path>
 - directory traversal
- STOR <file>
 - save file in the current directory on server
- RETR <file>
 - retrieve **file** from current directory on server
- PORT <B1, B2, B3, B4, B5, B6>
 - use TCP connection on port $B5 * 256 + B6$ on **B1.B2.B3.B4**

World Wide Web

- Goals
 - Allow browsers to browse hypertext documents stored on multiple servers



World Wide Web (2)

- The five key elements of [www](#)
 1. An addressing scheme that allows to identify any document stored on a server
 - **URL** : Uniform Resource Locator
 2. An hypertext language that allows to easily write documents with hypertext links
 - **HTML** : HyperText Markup Language
 3. An efficient and lightweight application-level protocol to exchange documents
 - **HTTP** : HyperText Transfer Protocol
 4. Servers
 5. Clients (browsers)

Uniform Resource Locator (URL)

□ Uniform Resource Locator (URL)

□ generic syntax : `<protocol>://<document>`

- `protocol` used to retrieve document from server
 - http is the most common one but others are frequently used
- `document` indicates the server and the location of the document
- `<user>:<password>@<server>:<port>/<path>`
 - `<user>` : optional username
 - `<password>` : optional password
 - `<machine>` : hostname or IP address of the server that hosts the document
 - `<port>` : optional port number
 - `<path>` : document location on server

□ examples

- <http://www.info.ucl.ac.be>
- <http://alice:secret@inl.info.ucl.ac.be:80/index.html>

HTML

□ HyperText Markup Language

- Language used to encode documents on the web

□ Keywords

- `<HTML>...</HTML>`
- `<HEAD>...</HEAD>`
- `<BODY>...</BODY>`
- `<TITLE>...</TITLE>`
- `...`
- `<I>...</I>`
- `<H1>...</H1>`
- `<P>`
- `<HR>`
- `...`
- `...`
- ``
- `text anchor`



HTML (2)

□ Example

```
<HTML>
```

```
<HEAD>
```

```
<TITLE>HTML test page</TITLE>
```

```
</HEAD>
```

```
<BODY>
```

```
<IMG SRC="http://www.images.be/logo.gif">
```

```
<H1>Web servers from UCL UCL</H1>
```

```
<HR>
```

```
<UL>
```

```
<LI><A HREF="http://www.uclouvain.be">UCL</A>
```

```
<LI><A HREF="http://www.info.ucl.ac.be">CSE Dept.</A>
```

```
<LI><A HREF="http://www.math.ucl.ac.be">Math</A>
```

```
</UL>
```

```
</BODY>
```

```
</HTML>
```

Header

Body

Image on remote server

First level title

External hypertext link

Information transfer www

□ HTTP 1.0 - non-persistent connection

□ Principle

- relies on TPC service (default port : 80)
- Client sends request, server sends reply

Client

Server




Information transfer [www](http://www.ietf.org/rfc/rfc2616.txt)

□ HTTP 1.0 - non-persistent connection

□ Principle

- relies on TPC service (default port : 80)
- Client sends request, server sends reply




Information transfer www

□ HTTP 1.0 - non-persistent connection

□ Principle

- relies on TPC service (default port : 80)
- Client sends request, server sends reply




Information transfer www

□ HTTP 1.0 - non-persistent connection

□ Principle

- relies on TPC service (default port : 80)
- Client sends request, server sends reply



HTTP

Client



Server



HTTP

Header contains additional information
about request sent by client

Method
GET
□ POST
□ ...

Request

Method
Header

CRLF

MIME Document

Client



Server



HTTP

Header contains additional information about request sent by client

Method
GET
□ POST
□ ...

Request




Client



Server



Response



Success or failure

Header contains information about server and optional parameters specific to response

HTTP

Header contains additional information about request sent by client

Method
GET
□ POST
□ ...

Request




Client



Server



Response



Success or failure

Header contains information about server and optional parameters specific to response

HTTP is a stateless protocol, server does not maintain any state from one request to another

POP, FTP, SMTP are examples of stateful protocols in contrast

HTTP : Example

Client




Server

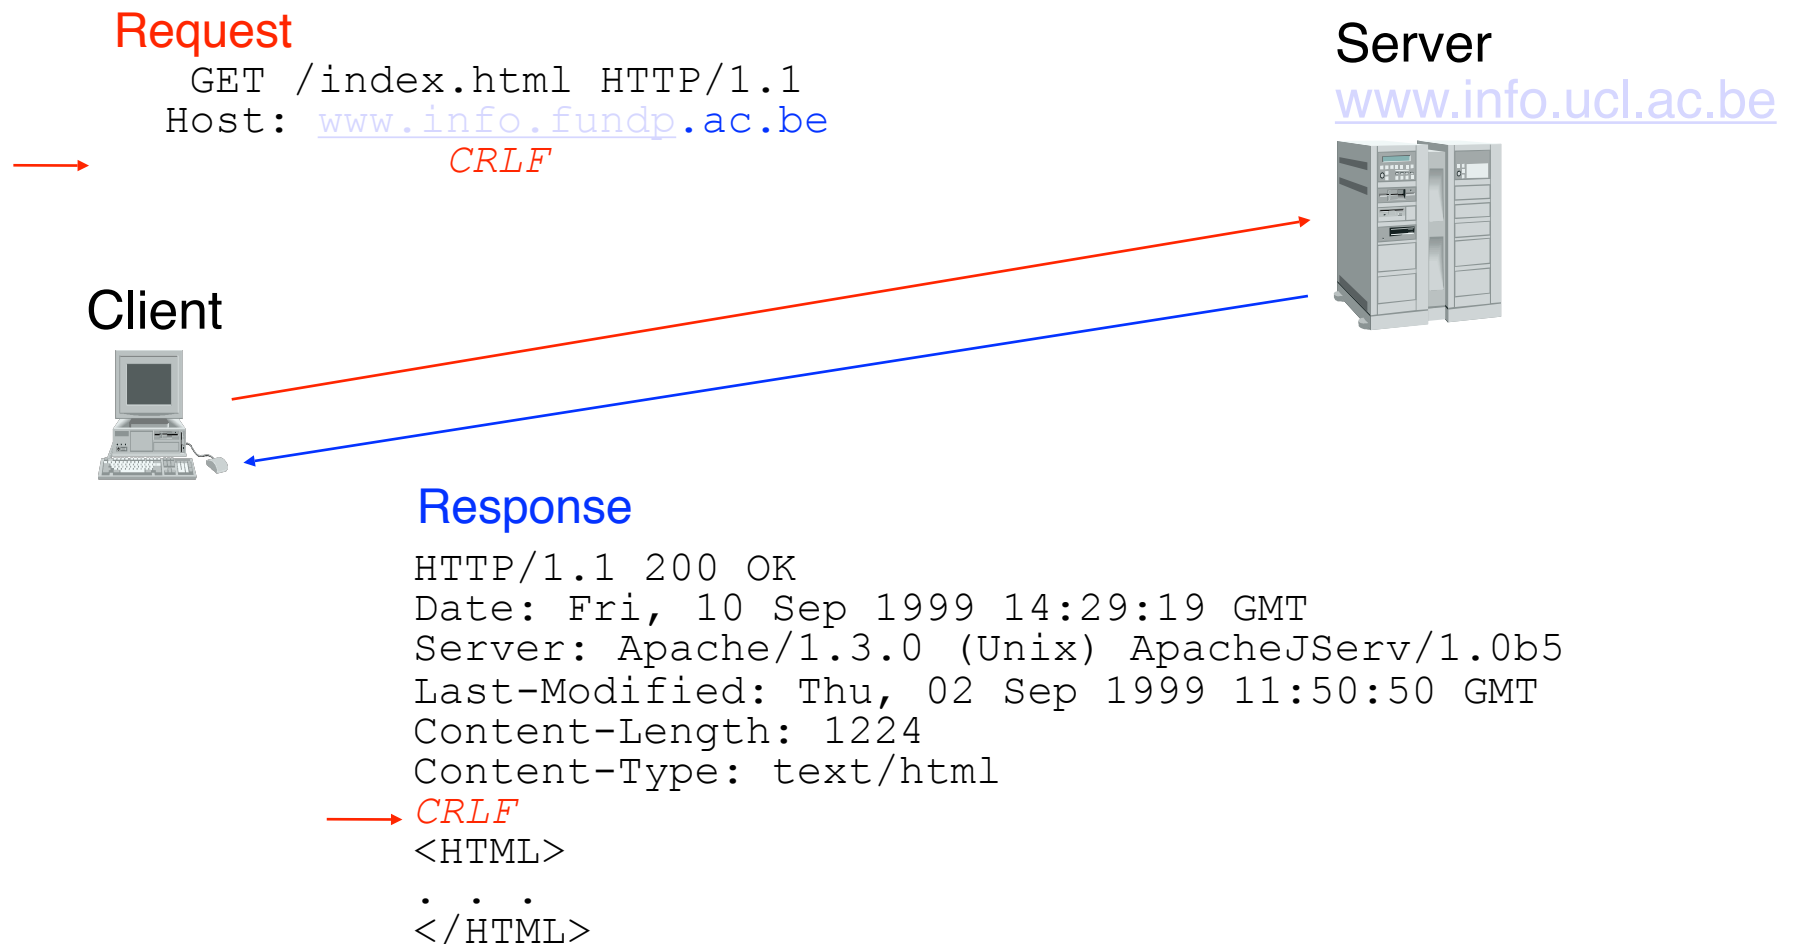
www.info.ucl.ac.be



HTTP : Example



HTTP : Example



HTTP : Methods

□ Methods

□ GET

- method used to request a "document" stored on server

- GET <document> HTTP/1.0

- example

- GET /index.html HTTP/1.0

□ POST

- method used to send a "document" to a server

- document is part of the request and encoded as a MIME document

HTTP : Request headers

□ Request headers

- Allow to add information about the client or the request

- Host: <name>

- Name of the server where the document is stored

- Authorization

- allows to perform access control

- If-Modified-Since: <date>

- server will only send the requested document if the document is more recent than date

- Referer: <url>

- Information, indicates the URL visited by the client before this request

- User-Agent: <agent>

- information, indicates the browser used on the client

HTTP : Status line

□ Status line

□ Format : Version_HTTP Code Comment

□ Success/Failure

□ 1xx : For information (unused)

□ 2xx : Success

□ Example : HTTP/1.0 200 OK

□ 3xx : Redirection

□ Request could not be handled on local server and should be sent to another server

□ Example :

□ HTTP/1.0 301 Moved permanently

□ attached MIME document will contain URL of document

□ 4xx : Client-side error

□ examples

□ syntax error, unreachable URL, unauthorised, ...

□ 5xx : Server-side error

□ examples :

□ internal error, method not implemented on server, ...

HTTP : Response headers

- Header

- Optional information about the server, the response or the document attached to the response

- Date

- date of the document attached to response
 - example : `Date: Wed, 05 Sep 2001 13:27:34 GMT`

- Server

- Name and version of http server used
 - example :
`Server: Apache/1.3.20 (Unix)ApacheJServ/1.1.2 PHP/4.0.6`

- Content-*

- MIME header of the attached document
 - example :
`Content-Length: 5891`
`Content-Type: text/html`

HTTP 1.1

HTTP 1.1

- HTTP 1.0
 - a single TCP connection used to transmit a single document (html file, image,...)
 - the establishment and release of the TCP connection induce a significant overhead, in particular for small pages

HTTP 1.1

□ HTTP 1.0

- a single TCP connection used to transmit a single document (html file, image,...)
 - the establishment and release of the TCP connection induce a significant overhead, in particular for small pages

□ HTTP 1.1

- uses a single persistent TCP connection
 - This TCP connection can be used for several requests and the corresponding responses
 - the cost of establishing and releasing the TCP connection is amortised over multiple requests
 - Although HTTP 1.1 uses a single TCP connection for multiple requests, HTTP 1.1 remains stateless

HTTP 1.1 : Persistent connection


Client




Server




HTTP 1.1 : Persistent connection



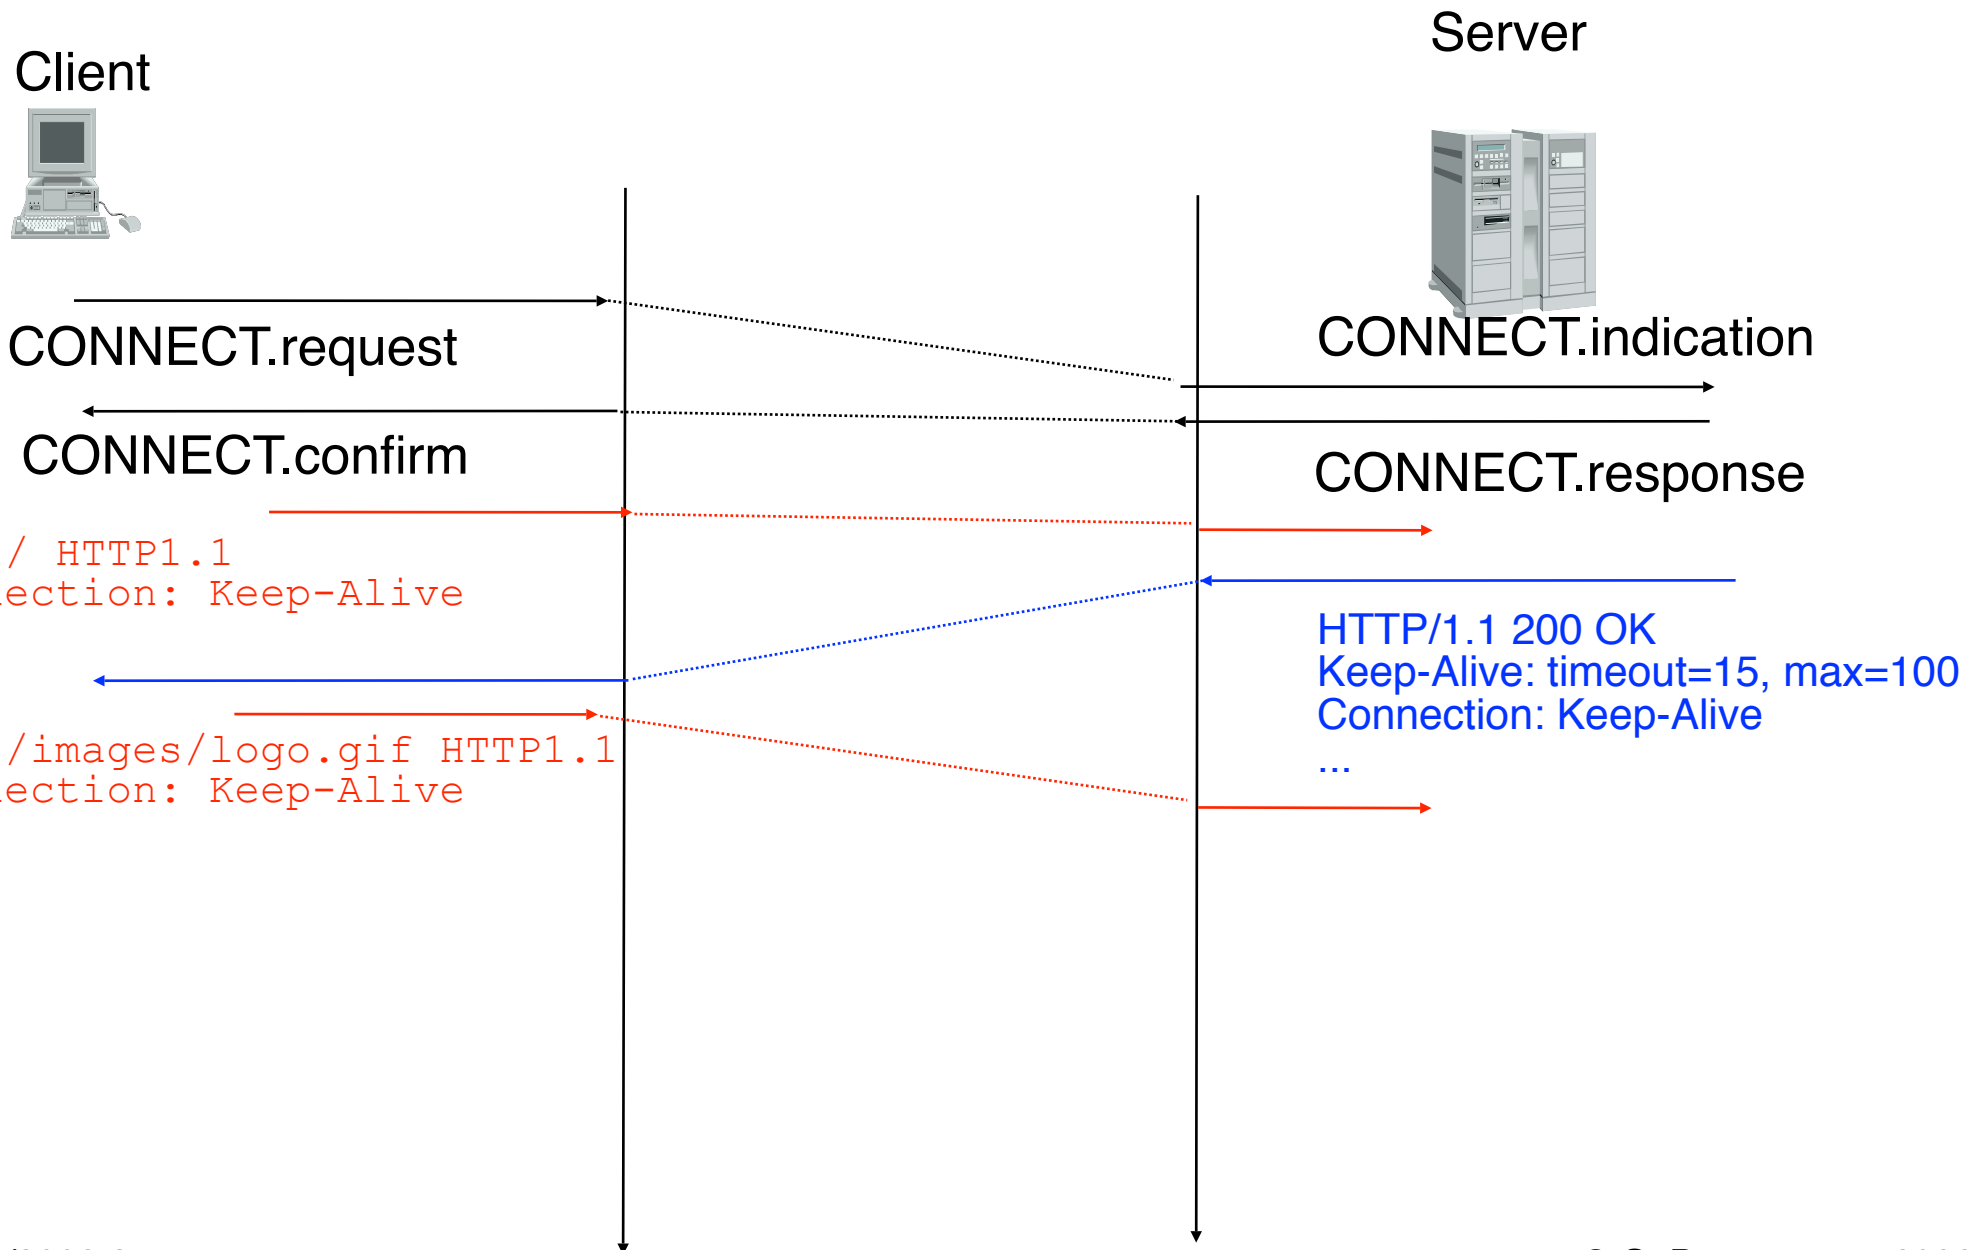
HTTP 1.1 : Persistent connection



HTTP 1.1 : Persistent connection



HTTP 1.1 : Persistent connection



HTTP 1.1 : Persistent connection

Client



Server



CONNECT.request

CONNECT.indication

CONNECT.confirm

CONNECT.response


GET / HTTP1.1
Connection: Keep-Alive
...

HTTP/1.1 200 OK
Keep-Alive: timeout=15, max=100
Connection: Keep-Alive
...

GET /images/logo.gif HTTP1.1
Connection: Keep-Alive
...

HTTP/1.1 200 OK
Keep-Alive: timeout=15, max=99
Connection: Keep-Alive
...

HTTP 1.1 : Persistent connection




Improving performance

- Observation
 - Many pages are requested multiple times or from close endhosts
- Solution
 - local cache on each client
 - `if-modified-since` header helps
 - one cache for multiple endhosts

Improving performance

- Observation
 - Many pages are requested multiple times or from close endhosts
- Solution
 - local cache on each client
 - if-modified-since header helps
 - one cache for multiple endhosts



HTTP Authentication

□ Example

Client



Server



HTTP Authentication

□ Example

Client



Server

GET / HTTP1.1
...

HTTP/1.0 401 Authorization req
WWW authenticate: machin
...

Browser asks user/password to user

HTTP Authentication

□ Example

Client



Server

GET / HTTP1.1

...

HTTP/1.0 401 Authorization req
WWW authenticate: machin

...

Browser asks user/password to user

GET / HTTP1.1

Authorization: User-password

...

HTTP/1.1 200 OK

...

GET /images/t.gif HTTP1.1

Authorization: User-password

...

Browser sends user/password in each request

HTTP Cookies


□ Example

Client




Server

Normal response



HTTP Cookies

□ Example



HTTP Cookies

□ Example

Client



Server

GET / HTTP1.1
...

HTTP/1.0 200 OK
Set-Cookie: machin
...

Normal response

Browser saves cookie

GET /doc HTTP1.1
Cookie: machin
...

HTTP/1.1 200 OK
...

*Response is function
of URL and cookie*

GET /images/t.gif HTTP1.1
Cookie: machi
...

*Browser sends cookie in all
requests sent to server*


Internet applications

- Contents


- The client-server model
- Name to address resolution
- email
- world wide web

→ □ peer-to-peer applications


Predicted IP traffic growth



Predicted IP traffic growth



Predicted IP traffic growth (3)



Peer-to-peer file sharing


- Evolution of file sharing on the Internet
- Servers using `ftp` protocol
 - A single server that serves all files on disk
 - A set of mirror servers serving the same content
- The innovation introduced by Napster
 - How to distribute many files from many nodes ?
 - Keep the files on their source nodes
 - Central Napster server stores description and URL of each shared file
 - Users willing to obtain a file consult central server to obtain file URL and then download file from their respective source nodes
 - server remains simple and can index large number of files
 - server does not directly participate in file transfer

Peer-to-peer file sharing (2)

- Limitations of the Napster approach
 - a single server indexes all files
 - if a source node fails, then the ongoing file transfers must be restarted
 - completely or partially depending on the file transfer protocol begin used for the transfer
 - performance of file transfer is function of performance of the corresponding source node
 - if source node is connected via ADSL, performance will be severely limited
- How to improve ?
 - Divide the file in blocks
 - Each block can be served by multiple nodes
 - provides redundancy
 - Download from several nodes at the same time
 - one TCP connection may be slow and others faster

Distributed Hash Table based P2P


- How to scale file sharing to a very large number n of nodes ?
- Principle of the solution
 - Use a hash function such as SHA-1
 - Each node has one identifier, $id = \text{hash}(\text{IP address})$ and a pointer to its successor on the Chord ring



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of `hash("FOOBAR")` on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle


- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to store files ?

□ Principle

- File FOOBAR is stored on the node whose id is the successor of $\text{hash}(\text{"FOOBAR"})$ on the ring
- A node on the ring uses its successors to find the responsible node for a given file
 - Examples



How to find files faster ?


How to find files faster ?

- Performance of file lookup
 - $O(n)$
 - worst case is to follow linked list of n nodes

How to find files faster ?

- Performance of file lookup
 - $O(n)$
 - worst case is to follow linked list of n nodes
- How to improve ?
 - Allow nodes to know additional pointers to other nodes on the Chord ring to speedup lookup
 - m = number of bits in the key/node identifiers
 - Each node maintains routing table of m entries
 - The i th entry in the table at node n contains the identity of the first node, s , that succeeds n by at least 2^{i-1} on the identifier circle, i.e., $s = \text{successor}(n + 2^{i-1})$
 - arithmetic modulo 2^m is used
 - A finger table entry includes both the Chord identifier and the IP address (and port number) of relevant node.

Scalable lookup with Chord




N8+1	N13
N8+2	N13
N8+4	N13
N8+8	N29
N8+16	N29
N8+32	N46

Scalable lookup with Chord

□ Example

□ $m=8$




$N8+1$	$N13$
$N8+2$	$N13$
$N8+4$	$N13$
$N8+8$	$N29$
$N8+16$	$N29$
$N8+32$	$N46$

Scalable lookup with Chord

□ Example

□ $m=8$




$N8+1$	N13
$N8+2$	N13
$N8+4$	N13
$N8+8$	N29
$N8+16$	N29
$N8+32$	N46

□ How to find key 53 from node 8 ?

Scalable lookup with Chord

□ Example

□ $m=8$



$N8+1$	$N13$
$N8+2$	$N13$
$N8+4$	$N13$
$N8+8$	$N29$
$N8+16$	$N29$
$N8+32$	$N46$


$N46+1$	$N51$
$N46+2$	$N51$
$N46+4$	$N51$
$N46+8$	$N57$
$N46+16$	$N57$
$N46+32$	$N8$

□ How to find key 53 from node 8 ?

Scalable lookup with Chord

□ Example

□ $m=8$



N8+1	N13
N8+2	N13
N8+4	N13
N8+8	N29
N8+16	N29
N8+32	N46


N46+1	N51
N46+2	N51
N46+4	N51
N46+8	N57
N46+16	N57
N46+32	N8

□ How to find key 53 from node 8 ?

Scalable lookup with Chord

□ Example

□ $m=8$



$N8+1$	$N13$
$N8+2$	$N13$
$N8+4$	$N13$
$N8+8$	$N29$
$N8+16$	$N29$
$N8+32$	$N46$

$N46+1$	$N51$
$N46+2$	$N51$
$N46+4$	$N51$
$N46+8$	$N57$
$N46+16$	$N57$
$N46+32$	$N8$

□ How to find key 53 from node 8 ?

Summary

- Client-server model
- UDP and TCP services
- Application-level protocols
 - DNS
 - relies on UDP, stateless
 - SMTP, POP, FTP
 - rely on TCP, statefull
 - HTTP
 - relies on TCP, stateless
- Peer-to-peer applications